


Why Does the CA SaaS Operating Platform Use OpenShift by Red Hat?

29.04.2016



The Problem

Let's consider an application with a back-end web service.




The service could be Tomcat serving HTML, Jetty serving OData, Node.js serving plain REST APIs, an instance of Ruby on Rails — doesn't matter for this argument.

How can we deploy this service?

The Problem

Let's say 1 instance of that service can handle **100 users** and requires **1 CPU core**.


And let's say we have 100 users. We could simply run 1 service on 1 physical server (or VM) with 1 core. No problem.



But what if we have **100,000** users?

The Problem

With 100,000 users, we need 1,000 service instances, which need at least 1,000 cores.



How can we deploy an application with **1,000 Tomcats** or **1,000 RoR instances**?

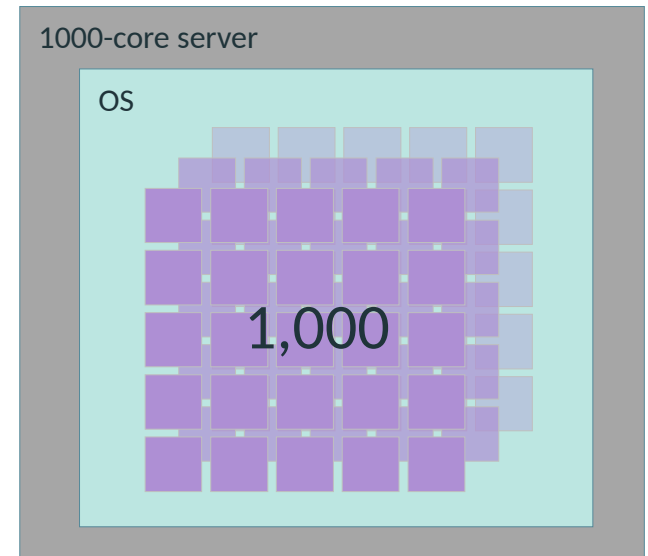
That's a complex problem, and good solutions must meet many criteria. But here are 4 key criteria for the CA SaaS Operating Platform.

Four Key Criteria

Capex	We're going large-scale in the first place because we are delivering SaaS, and SaaS profitability calls for the most cost-efficient hardware, software, and operations labor that can do the job.
Opex	
Isolation	To achieve sufficient reliability, we need resource isolation between service instances — we can't let one misbehaving instance take down hundreds of others.
Flexibility	We need to efficiently add, remove, grow, and shrink services. In real life, that hypothetical set of 1,000 service instances isn't static, isn't uniform, and isn't the only set of services.


Bad Approach: Humongous Servers

Putting all 1,000 services on a small number of very large servers is a non-starter. Biggest problem in general: Large servers are **dramatically less cost-effective** than small servers. Large servers only make sense for specialized workloads that can't run on small servers.



Capex	Awful
Opex	OK but doesn't matter
Isolation	Awful
Flexibility	Awful

Bad Approach: Horde of Tiny Servers




Putting each service on its own just-big-enough server doesn't work either. Even if that server size happens to be cost-efficient in cores-per-dollar, opex and flexibility are unacceptable.

Capex	Bad. Cost-optimal server size is usually bigger than 1 service, and some of our cost (such as OS licensing) scales with the number of servers.
Opex	Awful because we have 1,000 independent OS instances to manage. (The data center automation business from 10 years ago would love to help you try to fix this.)
Isolation	Good (actually, total overkill for web apps).
Flexibility	Awful especially if our service outgrows our server!


Common Approach: Virtualization

Let's say the cost-optimal server has 8 cores — enough raw capacity for 8 services (ignoring overhead). To isolate those 8 services from each other, we need some way to subdivide our physical servers. **Enter virtualization.**



Capex	Just OK, because those 1,000 OS instances impose significant additional overhead — we actually need way more than 125 servers.
Opex	Still awful because we still have 1,000 independent OS instances to manage. (Plus an additional ton of virtualization software to buy and operate — which is why VMware loves this approach.)
Isolation	Good (still overkill).
Flexibility	OK. We can change deployment sizing virtually, though not dynamically.

Aside: Intro to OS Container Technology



OS Container Technology

What if we could keep the flexibility of virtualization but remove the OS overhead?
Especially for mostly-uniform workloads that don't need total isolation?

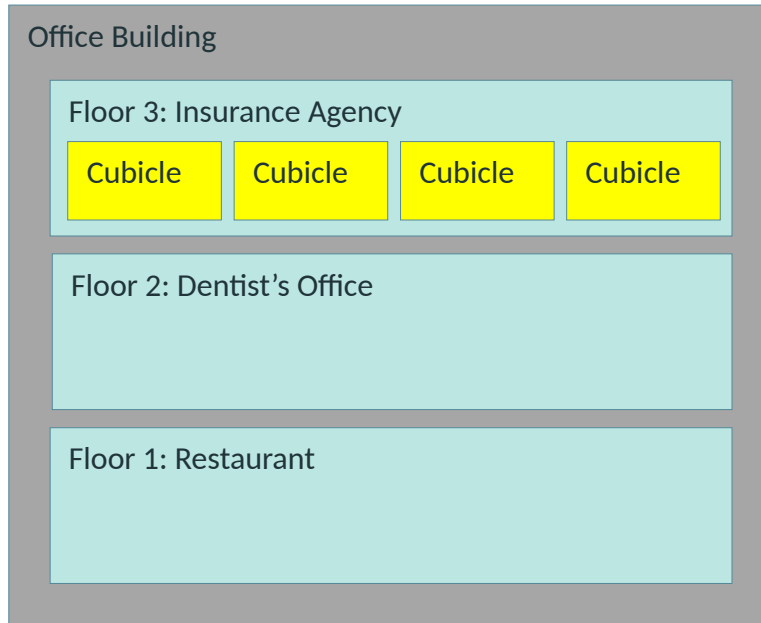
This observation led to the development of “lightweight virtualization” where the “VMs” are **partially isolated** partitions of the **same underlying OS**.

In full virtualization, one partition could be Windows, another Linux. With OS containers, both partitions expose the same OS instance.

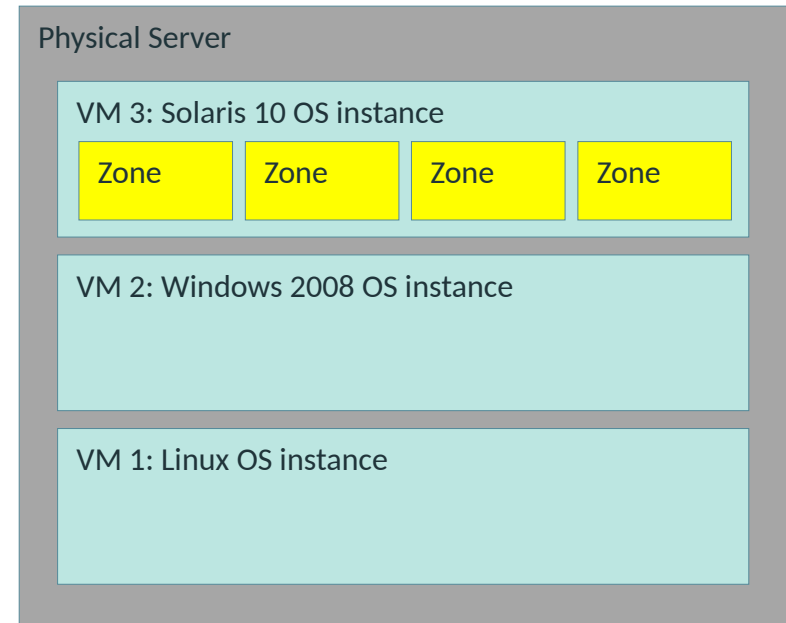
Virtualization is **external** to the guest OS. Containers are features provided **by** an OS.

There is a diverse spectrum of technologies for containerization, including Solaris Zones, Parallels Virtuozzo, Linux LXC containers, and BSD jails, ranging from low overhead/low isolation to high overhead/high isolation.

OS Container Analogy




Each floor is built out completely differently. The cubicles in the insurance agency are partially isolated — different office chairs, same carpet.




Each VM has a completely different OS. The containers in the Solaris instance are partially isolated — different process tables, same kernel.

Back to the Main Story



Common Approach: OS Containers (E.g., Joyent)



Capex	Good. Cost-optimal hardware. OS (and virtualization) overhead is at a minimum.
Opex	OK. Reasonable number of OS instances to manage, but we haven't addressed how to manage thousands of services or how to route application traffic. Typically that's a ton of custom automation.
Isolation	Good enough for web services.
Flexibility	OK. Containers are easier to manage than full VMs, but they are still individually-configurable objects that must be externally automated or manually maintained.

Toward a Solution

- With OS containers, capex and isolation are **good**.
- How can we extend that solution to fully-automated container management with sufficient flexibility for complex large-scale SaaS applications?

Introducing OpenShift by Red Hat

OpenShift is an open-source technology stack from Red Hat with several major features. For this discussion, we care about two big aspects of OpenShift:

- **Ultra-lightweight containers** — OpenShift combines two standard Linux features (cgroups and SELinux) to define lightly-isolated containers, called **gears**, that impose essentially zero overhead versus running on the bare OS. Gears provide just enough isolation for web services.

- **Distributed application model** — OpenShift allocates gears to applications from a pool of gears (running on a pool of OS instances). The number of gears assigned to an application can change dynamically in response to manual commands or automatic load-driven scaling decisions. And OpenShift automates a network of HTTP proxies so apps don't have to worry about which gears are running which services.

OpenShift-Based Approach to the Core Problem

Capex	Good. Cost-optimal hardware. OS (and virtualization) overhead is at a minimum.
Opex	Good. Reasonable number of uniform OS instances to manage, and OpenShift fully automates the containers, service deployment, and service traffic routing.
Isolation	Good enough for web services.
Flexibility	Good. All the dynamic large-scale complexity is software-defined and fully automated.

Why does the CA SaaS Operating Platform use OpenShift by RedHat? **Because it is the only scalable application framework that fully delivers on all 4 key criteria.**

How the CA SaaS Operating Platform Uses OpenShift



The CA SaaS Operating Platform leverages market-leading CA security products, such as Layer 7 and CloudMinder, for enterprise-grade authentication, security, and multitenancy.

Platform Security & Multitenancy
(Authentication/SSO, tenancy model, firewall, reverse proxy)

The CA SaaS Operating Platform provides a highly standardized model for its applications (CA SaaS products). Applications model business logic as platform **engines** (separately-scalable tiers) built on a small number of platform-provided application frameworks. The **CA SaaS Operating Platform is a PaaS that leverages OpenShift internally.**

Platform Application Model

Platform app 1 (CA SaaS product)

Engine 1 (Java OData API)


Engine 2 (Node.js REST API)

Platform app 2

Engine 1 (Java OData API)

The CA SaaS Operating Platform's automation uses the OpenShift API to run an OpenShift application for each engine. OpenShift allocates gears and dynamically manages a network of HTTP proxies.

OpenShift




OpenShift maintains a pool of gears on standardized RHEL OS instances (running on cost-efficient underlying IaaS or virtualized hardware — making CA's SaaS platform highly portable).

Platform Database-as-a-Service

MySQL

Cassandra



Q&A

